

ACID Atomicity
Consistency parallel connections
Isolation
Durability
 $RW \rightarrow Model \rightarrow Table$
model needs to model realworld

Entity must be uniquely identified by its attributes, not relation
ER diagram has entity sets, relation sets, aggregates etc.
— unconstrained
(A1₁ INT REFERENCES ENT₁(A1₁))

primary key is a combination so can have both
can have more than one value from each entity
zero participation
Ternary doctors, prescribes, drugs

—> at most 1 key constraint
(ENT1 INT PRIMARY KEY REFERENCES TABLE(ENT2))
< NULL, 1423 > is not necessary as it is already in the original table
how to know that a student is supervised that is why we have not null
give me all the students (without knowing if they are supervised) can use a
projection - sid - on the students table

— > with —
If ENT2 is not associated with any ENT1, it is not in . If it is associated
with ENT1, then PRIMARY KEY constraint ensures it appears only once.

— > with < —
A1₁ INT PRIMARY KEY REFERENCES ENT₁(A1₁), A1₂ INT UNIQUE
NOT NULL
REFERENCES ENT₂(A1₂),

=> exactly one total participation + key
each student must be supervised by exactly 1 prof.
Foreign key occurs where u can merge into a table
A1₁ INT **NOT NULL** REFERENCES ENT₁(A1₁),
A1₂ INT PRIMARY KEY REFERENCES ENT₂(A1₂),

== 1 or more

ISA

PRIMARY KEY REFERENCES Every student can be identified by Sid, then undergraduates can graduates must also be identifiable by Sid

Identity dependency

For a given city in a particular city

there can be multiple union city

combination to be a entity explains the references stateID

need to remove all the cities from indiana before removing indiana plus two addition delete statements

every identity dependency need an existential dependency

know the party exist

need a primary key to use a primary key

Overlapping constraint On delete Cascade (to avoid violate constraints)

When to use aggregation

FOREIGN KEY (A_{11}, A_{12}) REFERENCES $REL_1(A_{11}, A_{12})$

ENT1 = 20, ENT2 = 30 Minimum entries in relationship

ENT1 — <> — ENT2 0

ENT1 ==><>< — ENT2 20

ENT1 ==><> — ENT2 20

ENT1 —><>< — ENT2 0

Maximum entries in relationship

ENT1 — <> — ENT2 20 x 30 = 600

ENT1 ==><>< — ENT2 20 not enough to go around

ENT1 ==><> — ENT2 20

ENT1 —><>< — ENT2 20

nullary all attributes are primary key

unary association between an entity and its attributes

Sometimes, a constraint forces us to have circular references. In that case, there are several ways to solve this two of which are shown below:

Completely break the circularity and ignore certain constraints. If you wish, you can recover the circularity via ALTER TABLE.

horse owned by people to participate in a race

ISA

cannot be non covering

If there is one entity then it is redundant, need to capture part time does not need an office.

Joins

```
SELECT R.rname, R.area, S.price
FROM sells S NATURAL JOIN restaurants R
WHERE S.pizza = 'Funghi'
```

```
SELECT R.rname, (SELECT area from restaurants WHERE S.area = R.area)
, S.price
FROM sells
WHERE S.pizza = 'Funghi'
```

Aggregate

If DISTINCT is not specified, it will only look at Non-NULL values
WHERE will remove rows before the computation

```
SELECT R.rname
FROM restaurant
GROUP by R.rname
```

If you do the stable sorting yourself,
it is like sorting col2 DESC before col1 ASC

Get normalized mean

```
SELECT SUM(y) FROM (
  SELECT ((x - (SELECT MIN(x) + 0.0 FROM vals)) /
    (SELECT MAX(x) + 0.0 - MIN(x) + 0.0 FROM vals))
  AS y FROM vals GROUP BY x) AS z
```

is equivalent to

```
SELECT (SUM(x)-COUNT(x)*MIN(x))/(MAX(x)-MIN(x)+0.0)
```

Closed form formula = $\frac{sum_x(x_i)+sum_x(-min(X))}{max(X)-min(X)}$

Case analysis

```
SELECT * FROM players
ORDER BY
CASE WHEN title = 'Leader' THEN point END DESC,
CASE WHEN title = 'Minion' THEN point END ASC
which attribute to sort (title = ")
how to sort (ASC, DESC)
```

scalar subqueries $\leq 1row, = 1col, 0 \rightarrow NULL$

SQL EXCEPT

```
SELECT DISTINCT cname
FROM likes
EXCEPT ALL
SELECT DISTINCT cname
FROM likes
WHERE pizza IN (SELECT pizza FROM sells WHERE
rname = 'Corleone Corner');
```

is equivalent to

```
SELECT cname
FROM likes
EXCEPT
SELECT cname
FROM likes
WHERE pizza IN (SELECT pizza FROM sells WHERE
rname = 'Corleone Corner');
```

SQL EXISTS

```
SELECT DISTINCT cname
FROM likes
EXCEPT ALL
SELECT DISTINCT cname
FROM likes L
WHERE EXISTS (SELECT 1 FROM sells S WHERE
rname = 'Corleone Corner'
AND L.pizza = S.pizza);
```

SQL ANY

```
SELECT DISTINCT name
FROM sells
WHERE rname <> 'C'
AND price > ANY (SELECT price FROM sells WHERE rname = 'C')
```

is equivalent to

```
SELECT DISTINCT s1.rname
FROM sells s1, sells s2
WHERE s1.price > s2.price AND s2.rname = 'Corleone Corner'
AND s1.rname <> 'Corleone Corner'
```

Total order property

```
SELECT DISTINCT name
FROM sells
WHERE rname <> 'C'
AND price > (SELECT MIN(price) FROM sells WHERE rname = 'C')
```

SQL ALL

$$\pi_{rname, pizza, price}(\sigma_{s1price > s2price}(\rho_{(s1price, rname)}(sells) \times \sigma_{(s2price, rname)}(sells)))$$

```

SELECT DISTINCT rname, pizza, price
FROM sells S1
WHERE S1.price >= ALL (SELECT S2.price
FROM sells S2
WHERE S1.rname = S2.rname)

```

Scoping

ambiguous

```

SELECT * FROM (SELECT * FROM sells) AS T, (SELECT * FROM
T) AS R

```

CTE

```

WITH intermediate AS
(SELECT C.cid, C.name, COUNT(*)
FROM enrolls E NATURAL JOIN courses C
GROUP BY C.cid)

SELECT name, num
FROM intermediate
WHERE num > (SELECT num FROM intermediate
            WHERE name = "Database Systems"
            AND cid = 'CS2102');

```

```

WITH X AS ... WHERE a > 20
     Y AS ... WHERE b < 10

```

E.g. a not present cte

```

WITH
not_present AS (
SELECT sid, week
FROM (SELECT sid FROM Students) AS tbl,
(SELECT DISTINCT week FROM presenters) AS wk_tbl

```

Where instead of having

```
WITH CTE
SELECT COUNT(*)
```

```
SELECT FROM CTE
WHERE
```

Consecutive values in a table

Suppose $n=3$

q1 = not in table and $\max(n) \leq 3$

q2 = 4th onwards

q3 = $\max - n$

q4 = between n and $n - 3 - 1$

Variables

capturing a computation process CTE: common table expression (temporary table unlike VIEW which is persistent)

imperative, computes line by line, don't care how it is computed

E.g. `SELECT * FROM __ WHERE x is parent and x is male`

Likes all pizza sold by a restaurant

If restaurants has pizza A and B
customers likes A, B, C
this subset should be still be included

Coalesce

If there is no entry or is null, fill with 0

```
SELECT COALESCE(MIN(num), 0) FROM status;
```

Universal Quantification

Intuition

$|P_1| \geq |P_2|$ for $\{1, 2, 3\} \not\subseteq \{4, 5\}$ $S1 \cup S = S1$ when S1 already contains S
 $S1 \cap S = S$ as they are incomparable

Double negation There is no pizza sold by CC where not (P likes pizza)

Functional Dependencies

Reflexivity

$AB \rightarrow A$

Composition

$A \rightarrow B$

$C \rightarrow D$

$AC \rightarrow BD$

Logical implication

Does it hold for all instances in the table

If FD is R(A, B, C, D)
then total number of FDs is $(2^4 - 1) \times (2^4 - 1)$
excluding itself so -1

$A \rightarrow BC$ is also $A \rightarrow B$ and $A \rightarrow C$

Attribute redundancy

$AB \rightarrow C$ and $B \rightarrow A$

We can replace $AB \rightarrow C$ with $A \rightarrow C$

FD redundancy - remove without changing its meaning

$A \rightarrow B$ and $B \rightarrow C$ means $A \rightarrow C$ is redundant

For each FD, check for every single node on the destination and if there is a path from source to destination

minimum cover is not unique

How to get key

If there is no arrow pointing to them in the hyper-graph, we need to include it as part of the key

How to get attribute closure

For e.g. $AD^+ = ACDE$

Using the hyper-graph connect the source to the nodes at the left side of the implication

- Is there a path that can reach the nodes given at the right side of the implication

How to get minimal cover

Usually the implications that can't be decomposed

Cannot have redundant attributes

$AB \rightarrow CDE$

B^+ check if B can determine CDE

which means A is redundant

$B \rightarrow CDE$

How to get key

Given a superkey CDG+
 Then try all combinations
 CD+ CG+ DG+
 Find attributes that is missing from the rhs

Decomposition


Intuition is like splitting the relation into two or more

Lossless join

If the relation only contain attributes on RHS most likely it is not lossless

$F \models (R1 \cap R2)$ satisfy a FD

1. Combine two relations into one
2. Each R can only be used one in the tree



Dependency preserving

Cannot include anything that is trivial

As long it satisfy one of the attribute on LHS, then the FD is satisfied.
 Can use the output to another FD.

Dependency-preserving

$F^+ = \{a \rightarrow b \mid F \models a \rightarrow b\} \quad F_a \equiv E$

Functional Dependency Projection

- $F[a] = \{b \rightarrow c \in F^+ \mid (b, c) \in a\}$
- Let $R(A, B, C)$ with $F = \{A \rightarrow B, B \rightarrow C, C \rightarrow B\}$
- Compute $F[AC] = \emptyset$ ~~Wrong~~

$A \rightarrow C \in F^+$

$F[AC] = \{A \rightarrow C\}$

$b \in a$

$b \rightarrow b^+$

$b \rightarrow b \cup a$

$\in F[a]$

Does R1 contain the attributes that is described by any FD?

Dependency-preserving

Functional Dependency Projection

- Let $R(A, B, C)$ with $F = \{A \rightarrow B, A \rightarrow C\}$
- Is $\{R_1(A, B), R_2(B, C)\}$ dependency-preserving?

$F[R_1] = \{A \rightarrow B\}$

$F[R_2] = \{\}$

$F[R_1] \cup F[R_2] \neq F$