

Some terms

Collaborative filtering - collaborative filtering is a method of making automatic predictions (filtering) about the interests of a user by collecting preferences or taste information from many users

Assumes that one person has same opinion as the other person

Inductive learning Assumption - after seeing many training examples result in high accuracy on unobserved examples

Hypothesis approximates well

average hypothesis only depends on H , does not always lie in H

$$H_2 \subset H_{10}$$

Target function $y = f(x)$

Unbiased learner cannot generalise \rightarrow inductive bias

h is +ve when $h(x) = 1 \geq_g$ more general than or equal to

$$h_j \geq_g h_k \equiv \forall x \in X (h_k(x) = 1) \rightarrow (h_j(x) = 1)$$

Not a total order as there may have the same specificity but different hypothesis (maximally specific h)

Consistency: h is consistent iff $h(x) = f(x)$

$$h4(\langle 0, 1 \rangle) = 1 \neq f(\langle 0, 1 \rangle) = 0.$$

Proposition 1

h is consistent with iff every +ve training instance satisfies h and every -ve training instance does not satisfy h .

Proposition 2

If h belongs to H , then it is consistent

Naive Bayes

Bayes Rule

There can be relatively lesser occurrences of one class

prior odds / likelihood = posterior odds

Bayesian thinking

How likely is something not going to work?

Hold your beliefs, shift your beliefs as u encounter with the world

$$P(y|x) = \frac{P(x|y) \times P(y)}{P(x)}$$

$$P(X_1 = 0, X_2 = -1) = P(X_1 = 0 \mid C_1) \times P(X_2 = -1 \mid C_1) \times P(C_1) + P(X_1 = 0 \mid C_2) \times P(X_2 = -1 \mid C_2) \times P(C_2)$$

a priori - equally probable

Naive - Assuming that features are conditionally independent of each other

Insensitive to the number of training examples

Calculate which is more likely

1. $P(Outcome|x) = P(X_1|Outcome) * P(X_2|Outcome) * P(Outcome)$
2. $P(Not Outcome|x) = P(X_1|Not Outcome) * P(X_2|Not Outcome) * P(Not Outcome)$

Calculate probability outcome occurs

$$\#1 / (\#1 + \#2)$$

Perceptron classifier

$$x_0\theta_0 + x_1\theta_1 + x_2\theta_2$$

of the previous weight vector

first instance will compute based on initial vector

For -1, first point checks == -1, != 1 for the points after

Update weight = - (x × learning_rate) for the misclassified instance

For +1, first point checks != -1, == 1 for the points after

Update weight = + (x × learning_rate) for the misclassified instance

points are linearly separable if there are no longer any misclassified points

not misclassified: sign(x) function having -1 when x is ≤ 0

Supervised learning

Identifying objects

Output would be a continuous value

Given many emails, you want to **determine** if they are Spam or Non-Spam emails.

Given historical weather records, **predict** if tomorrow's weather will be sunny or rainy.

Unsupervised learning

No labels (no output) Kmeans

Input: Given a set of data points P , number of clusters k

1. Randomly pick k points from P as centers $c_j = 1..k$
 2. Iterate (until max-iterations or c_j no longer changes up to some threshold)
- Assign each point to nearest center: $y_i = \arg \min \| p_i - c_j \|^2$
 - c_j = Re-estimate each center as mean of points assigned to it

Given a set of news articles from many different news websites, find out what are the main topics covered.

From the user usage patterns on a website, figure out what different groups of users exist.

RL

Key usage: Planning problems

Goal of markov property: operate in ease
choices that maximise the outcome

Precision vs control

Decisions we make the optimize

Prediction allows us to evaluate how good a policy is for a state space

Control

Reward at any step is optimal, highest reward

continuous MDP - common in robots

No supervisor, only reward signal

Reward is a scalar feedback

whether one state vs the other is preferred

RL Challenge: Does not have a good idea if it is right to do, which one

is feasible, figuring out rules, which gives the best reward
Given action will give observation and reward
Does not know the rules

Planning Challenge:

A* search - Going through the search space to figure the appropriate action
what if i did this, what kind of consequences

Rules are known

Uses a search space which requires tree search

Action affects the subsequent states

Action \rightarrow Observation (what state it is in) \rightarrow Reward

time t will not take the final reward but $t - 1$

Sometimes actions affect the observations

May contain irrelevant information of the state at time t

Although we have large field of vision, our neural cortex is trying to observe
 t changes in the environment

Function of history is not visible to the environment and is specific to agent

State - what information that we have in the brain is what we are going to
need to predict the future

Fully observable

The state can be represented different based on how we want the agent to
see the sequence

congruent between the agent and the environment

Partially observable

Align/ localize itself based on the map e.g. game

Represented state can influence the agent

Depending on the agent, we can end up in a different state

E.g. if we want to turn the rotor around the corner, it may turn out not to
behave that way

Recurrent neural networks

Using ReLu to construct each state in deep learning

Policy

Look through the mapping function

Deterministic s_1 to a_1

Stochastic Take a particular action with a probability

Value function - future reward

Sum of expected reward, reward farther in the future will have gamma to discount it

Model - probability that it will go to the next state, sum over all possible actions has to be 1

Agent taxonomy

Value Based

Policy - in this state we take a particular action

Actor Critic - model based

Exploration vs Exploitation

Going to the restaurant that i like

Exploration - discover more information about the environment

Exploitation - Exploits known information to maximise reward

State transition probability matrix

Represented with the symbol P

some of the values are 0 in the transition matrix if it cannot move to a particular state

episodes are finite

there are always probability attach to the state, eventually it will reach the terminal state

0.1 Markov Reward process

Attach the reward and discount

Differentiating reward, create two different states with different reward

$E(r_{t+1}|..)$ immediately receive the reward in next time step

Discount factor γ

Near sighted - given where i am now where should I go

Far sighted - sum all the rewards add all of them, favor longer samples (should have more reward)

Role of Discount

Infinitely sample the cycle, a lot of uncertainties on the future

0.2 Bellman equation

Reward based on executing the state

Find S

1. Initialise h to most specific h
2. If the positive instance differs, then replace with ?

Version space

List all hypothesis, remove any hypothesis that is inconsistent with any training example

Candidate Elimination

prefers positive examples over negatives as there can only be one maximally specific hypothesis in each iteration ? don't care

no value

Set $S_0 = \langle \emptyset \emptyset \emptyset \emptyset \emptyset \emptyset \rangle$

$G_0, G_1 = \langle \text{?????} \rangle$

For the negative example,

Create h from the specific hypothesis when it differs based on previous instance

Remove those that are inconsistent from the previous general hypotheses

Keep S to the previous instance S

For the positive example,

replace with ? when it differs

Cost functions

Logistic regression

Is binary classification

Need meta strategies for multi class and it chooses the model that has the highest confidence in predictions

Only gives values 0 and 1 so it is not a good result

$$J(\theta) = \frac{1}{m} \sum_{j=1}^m \ln(1 + \exp^{-y^j \theta^T x^j})$$

Linear regression

$$J(\theta) = \frac{1}{m} \sum_{j=1}^m (\theta^T x^j - y^j)^2$$

Linear algebra

$$\text{trace}(\text{AB}) = \text{trace}(\text{BA}) = \text{trace}(\text{I}_{n+1}) = n + 1$$

Idempotency

$$H^2 = H$$

$$(I - H)^2 = I - 2H + H = I - H$$

Pseudoinverse

$$(X^T X)^{-1} X^T$$

$$\theta = X^+ y$$

$$X = \begin{pmatrix} 1 & .. & .. \\ 1 & .. & .. \\ 1 & .. & .. \end{pmatrix}$$

Maximizing likelihood

$$\text{Minimize cross entropy} = - \frac{1}{m} \sum_{j=1}^m \ln g(y^j \theta^T x^j)$$

Ridge Regression

$$\frac{1}{m} \sum_{i=1}^m (y_i - w^T x_i)^2 + \lambda \|w\|^2$$

$$-2X^T(Y - X * \theta) + 2 * \alpha * \theta$$

Noise

Not deterministic only deterministic for $f(x) = 0$

$$y - f(x)$$

$$P(y|x)$$

$$\text{Noisy target} = E[y|x] + (y - f(x))$$

Stochastic noise vs Deterministic

To reduce stochastic noise, the only way is to re-measure y

To reduce deterministic noise, change H

Underfitting and Overfitting

Why do we regularize? We fit the data too much, we are fitting to the noise, balance fitting to the observed data. If we fit the data too well, the performance can be poor called overfitting. Overfitting can happen even if the data is not noisy.

Overfit - high level of noise, high complexity of $f(x)$

Underfit - high bias and low variance

Reduce bias by increasing complexity of hypotheses
simple representation of more complex reality

Non-linear transformation

Transform this non-linear equation to a linear one: $y = ax^b$?

Under the condition $y > 0$ and $x > 0$

$$\ln(y) = \ln(ax^b)$$

$$\ln(y) = \ln(a) + \ln(x^b)$$

$$\ln(y) = \ln(a) + b\ln(x)$$

let $y' = \ln(y)$, $x' = \ln(x)$, $a' = \ln(a)$, the equation can be written as
 $y' = a' + bx'$