

## Some terms

Collaborative filtering - collaborative filtering is a method of making automatic predictions (filtering) about the interests of a user by collecting preferences or taste information from many users

Assumes that one person has same opinion as the other person

Inductive learning Assumption - after seeing many training examples result in high accuracy on unobserved examples

Hypothesis approximates well

average hypothesis only depends on  $H$ , does not always lie in  $H$

$$H_2 \subset H_{10}$$

Target function  $y = f(x)$

Unbiased learner cannot generalise  $\rightarrow$  inductive bias

$h$  is +ve when  $h(x) = 1 \geq_g$  more general than or equal to

$$h_j \geq_g h_k \equiv \forall x \in X (h_k(x) = 1) \rightarrow (h_j(x) = 1)$$

Not a total order as there may have the same specificity but different hypothesis (maximally specific  $h$ )

Consistency:  $h$  is consistent iff  $h(x) = f(x)$

$$h4(\langle 0, 1 \rangle) = 1 \neq f(\langle 0, 1 \rangle) = 0.$$

### Proposition 1

$h$  is consistent with iff every +ve training instance satisfies  $h$  and every -ve training instance does not satisfy  $h$ .

### Proposition 2

If  $h$  belongs to  $H$ , then it is consistent

## Naive Bayes

### Bayes Rule

There can be relatively lesser occurrences of one class

prior odds / likelihood = posterior odds

Bayesian thinking

How likely is something not going to work?

Hold your beliefs, shift your beliefs as u encounter with the world

$$P(y|x) = \frac{P(x|y) \times P(y)}{P(x)}$$

$$P(X_1 = 0, X_2 = -1) = P(X_1 = 0 \mid C_1) \times P(X_2 = -1 \mid C_1) \times P(C_1) + P(X_1 = 0 \mid C_2) \times P(X_2 = -1 \mid C_2) \times P(C_2)$$

**a priori** - equally probable

Naive - Assuming that features are conditionally independent of each other

Insensitive to the number of training examples

Calculate which is more likely

1.  $P(Outcome|x) = P(X_1|Outcome) * P(X_2|Outcome) * P(Outcome)$
2.  $P(Not Outcome|x) = P(X_1|Not Outcome) * P(X_2|Not Outcome) * P(Not Outcome)$

Calculate probability outcome occurs

$$\#1 / (\#1 + \#2)$$

## Perceptron classifier

$$x_0\theta_0 + x_1\theta_1 + x_2\theta_2$$

of the previous weight vector

first instance will compute based on initial vector

For -1, first point checks == -1, != 1 for the points after

Update weight = - (x × learning\_rate) for the misclassified instance

For +1, first point checks != -1, == 1 for the points after

Update weight = + (x × learning\_rate) for the misclassified instance

points are linearly separable if there are no longer any misclassified points

not misclassified: sign(x) function having -1 when x is  $\leq 0$

## Supervised learning

Identifying objects

Output would be a continuous value

Given many emails, you want to **determine** if they are Spam or Non-Spam emails.

Given historical weather records, **predict** if tomorrow's weather will be sunny or rainy.

## Unsupervised learning

No labels (no output) Kmeans

Input: Given a set of data points  $P$ , number of clusters  $k$

1. Randomly pick  $k$  points from  $P$  as centers  $c_j = 1..k$
  2. Iterate (until max-iterations or  $c_j$  no longer changes up to some threshold)
- Assign each point to nearest center:  $y_i = \arg \min \| p_i - c_j \|^2$
  - $c_j$  = Re-estimate each center as mean of points assigned to it

Given a set of news articles from many different news websites, find out what are the main topics covered.

From the user usage patterns on a website, figure out what different groups of users exist.

## RL

Key usage: Planning problems, app placements product recommendations

### Goal

Low risk when performing actions

Goal of markov property: operate in ease  
choices that maximise the outcome

Precision vs control

Decisions we make the optimize

Prediction allows us to evaluate how good a policy is for a state space

Control

Reward at any step is optimal, highest reward

continuous MDP - common in robots

No supervisor, only reward signal

Reward is a scalar feedback

whether one state vs the other is preferred

RL Challenge: Does not have a good idea if it is right to do, which one is feasible, figuring out rules, which gives the best reward  
Given action will give observation and reward  
Does not know the rules

Planning Challenge:

A\* search - Going through the search space to figure the appropriate action  
what if i did this, what kind of consequences  
Rules are known  
Uses a search space which requires tree search

Action affects the subsequent states  
Action  $\rightarrow$  Observation (what state it is in)  $\rightarrow$  Reward  
time  $t$  will not take the final reward but  $t - 1$   
Sometimes actions affect the observations  
May contain irrelevant information of the state at time  $t$   
Although we have large field of vision, our neural cortex is trying to observe  $t$  changes in the environment  
Function of history is not visible to the environment and is specific to agent  
State - what information that we have in the brain is what we are going to need to predict the future

### **Fully observable**

The state can be represented different based on how we want the agent to see the sequence  
congruent between the agent and the environment

### **Partially observable**

Align/ localize itself based on the map e.g. game  
Represented state can influence the agent  
Depending on the agent, we can end up in a different state  
E.g. if we want to turn the rotor around the corner, it may turn out not to behave that way  
Recurrent neural networks  
Using ReLU to construct each state in deep learning

## **Policy**

Look through the mapping function

Deterministic  $s_1$  to  $a_1$

Stochastic Take a particular action with a probability

Value function - future reward

Sum of expected reward, reward farther in the future will have gamma to discount it

Model - probability that it will go to the next state, sum over all possible actions has to be 1

## **Agent taxonomy**

Value Based

Policy - in this state we take a particular action

Actor Critic - model based

## **Exploration vs Exploitation**

Going to the restaurant that i like

Exploration - discover more information about the environment

Exploitation - Exploits known information to maximise reward

## **State transition probability matrix**

Represented with the symbol  $P$

some of the values are 0 in the transition matrix if it cannot move to a particular state

episodes are finite

there are always probability attach to the state, eventually it will reach the terminal state

### **0.1 Markov Reward process**

Attach the reward and discount

Differentiating reward, create two different states with different reward

$E(r_{t+1}|..)$  immediately receive the reward in next time step

Discount factor  $\gamma$

Near sighted - given where i am now where should I go

Far sighted - sum all the rewards add all of them, favor longer samples  
(should have more reward)

### **Role of Discount**

Infinitely sample the cycle, a lot of uncertainties on the future

## **0.2 Bellman equation**

Reward based on executing the state

### **Value iteration**

Idea of value iteration

Instead of looking what policy to take, look back and determine what is the best action brought us to get the best value function

Emits the observation and reward

Recover information or  $t$  and  $r$

Value iteration is a simple form of policy iteration

Use optimal  $V^*$  to create back the policy

### **Q learning**

Use  $q$  value to engineer the policy (monodically increasing)

$r(s, a)$  reward at comes at next state

Higher value of  $q$  means we see it a lot of times, smooth it over time

## Neural Networks

### Motivation

In traditional ml, it is difficult to handcraft features, neural networks are good in feature learning

NN generalise badly

As the number of parameters (weights and layers) increase in a neural network, it can easily overfit the data.

Backpropagation is a complicated and tedious process. Moreover, the surface plotted by the error of a neural network in non-convex and high-dimensional and can consist of many hills and valleys - making it hard to reach the optimal set of weights.

Do all the neurons in a layer always share the same bias value?

Yes, that is correct. The bias value is the same value ( $x_0 = 1$ ) for all units at all layers, because we are going to learn the particular weight for the bias for each individual unit,  $\theta_0$ , separately. It's just pictorially cleaner to think that each layer has a separate bias unit, but it doesn't matter, because the value of the bias unit is always (by convention) set to unity, such that the individual weights  $\theta_0$  (or in the NN matrix notation  $\theta_0^{(l)}$ ) are the controls for the biases.

Random initialisation of weights for sign function

the  $\text{sign}()$  function takes different constant values in different regions and therefore has gradient 0.

the  $\text{sign}()$  function is discontinuous at  $x=0$ .

Stack in together in a linear unit

Generalise things that are out of sample

Regularisation for NN

Weight decay square of the weights are small so it is more toward linear regression

Is the Soft Weight Elimination function a Cost Function?

No, but it closely related. It is a regulariser, a separate, artificial part of

the augmented loss function, as what we learned in Week 05. We denote regulariser functions as  $\Omega(\theta)$ .

Would it have been a good idea to stop the first time the validation error goes up?

The validation error might fluctuate and therefore stopping it the first time it goes up will not give the optimal result. One can consider using a moving average or some other metric.

## Find S

1. Initialise  $h$  to most specific  $h$
2. If the positive instance differs, then replace with ?

## Version space

List all hypothesis, remove any hypothesis that is inconsistent with any training example

## Candidate Elimination

prefers positive examples over negatives as there can only be one maximally specific hypothesis in each iteration ? don't care

no value

Set  $S_0 = \langle \emptyset \emptyset \emptyset \emptyset \emptyset \emptyset \rangle$

$G_0, G_1 = \langle \text{??????} \rangle$

For the negative example,

Create  $h$  from the specific hypothesis when it differs based on previous instance

Remove those that are inconsistent from the previous general hypotheses

Keep  $S$  to the previous instance  $S$

For the positive example,

replace with ? when it differs



## Cost functions

Logistic regression

Is binary classification

Need meta strategies for multi class and it chooses the model that has the highest confidence in predictions

Only gives values 0 and 1 so it is not a good result

$$J(\theta) = \frac{1}{m} \sum_{j=1}^m \ln(1 + \exp^{-y^j \theta^T x^j})$$

Linear regression

$$J(\theta) = \frac{1}{m} \sum_{j=1}^m (\theta^T x^j - y^j)^2$$

## Linear algebra

$$\text{trace}(\mathbf{AB}) = \text{trace}(\mathbf{BA}) = \text{trace}(\mathbf{I}_{n+1}) = n + 1$$

*Idempotency*

$$H^2 = H$$

$$(I - H)^2 = I - 2H + H = I - H$$

## Pseudoinverse

$$(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$$

$$\theta = \mathbf{X}^+ y$$

$$\mathbf{X} = \begin{pmatrix} 1 & \dots & \dots \\ 1 & \dots & \dots \\ 1 & \dots & \dots \end{pmatrix}$$

## Maximizing likelihood

$$\text{Minimize cross entropy} = - \frac{1}{m} \sum_{j=1}^m \ln g(y^j \theta^T x^j)$$

## Ridge Regression

$$\frac{1}{m} \sum_{i=1}^m (y_i - w^T x_i)^2 + \lambda \|w\|^2$$
$$-2X^T(Y - X * \theta) + 2 * \alpha * \theta$$

## Noise

Not deterministic only deterministic for  $f(x) = 0$

$y - f(x)$

$P(y|x)$

Noisy target =  $E[y|x] + (y - f(x))$

## Stochastic noise vs Deterministic

To reduce stochastic noise, the only way is to re-measure  $y$

To reduce deterministic noise, change  $H$

## Underfitting and Overfitting

Why do we regularize? We fit the data too much, we are fitting to the noise, balance fitting to the observed data. If we fit the data too well, the performance can be poor called overfitting. Overfitting can happen even if the data is not noisy.

**Overfit** - high level of noise, high complexity of  $f(x)$

**Underfit** - high bias and low variance

Reduce bias by increasing complexity of hypotheses

simple representation of more complex reality

## Non-linear transformation

Transform this non-linear equation to a linear one:  $y = ax^b$ ?

Under the condition  $y > 0$  and  $x > 0$

$$\ln(y) = \ln(ax^b)$$

$$\ln(y) = \ln(a) + \ln(x^b)$$

$$\ln(y) = \ln(a) + b\ln(x)$$

let  $y' = \ln(y)$ ,  $x' = \ln(x)$ ,  $a' = \ln(a)$ , the equation can be written as  
 $y' = a' + bx'$