

Some terms

Collaborative filtering - collaborative filtering is a method of making automatic predictions (filtering) about the interests of a user by collecting preferences or taste information from many users

Assumes that one person has same opinion as the other person

Inductive learning Assumption - after seeing many training examples result in high accuracy on unobserved examples

Hypothesis approximates well

average hypothesis only depends on H , does not always lie in H

$$H_2 \subset H_{10}$$

Target function $y = f(x)$

Unbiased learner cannot generalise \rightarrow inductive bias

h is +ve when $h(x) = 1 \geq_g$ more general than or equal to

$$h_j \geq_g h_k \equiv \forall x \in X (h_k(x) = 1) \rightarrow (h_j(x) = 1)$$

Not a total order as there may have the same specificity but different hypothesis (maximally specific h)

Consistency: h is consistent iff $h(x) = f(x)$

$$h4(\langle 0, 1 \rangle) = 1 \neq f(\langle 0, 1 \rangle) = 0.$$

Proposition 1

h is consistent with iff every +ve training instance satisfies h and every -ve training instance does not satisfy h .

Proposition 2

If h belongs to H , then it is consistent

Naive Bayes

Bayes Rule

$$P(y|x) = \frac{P(x|y) \times P(y)}{P(x)}$$

$$P(X_1 = 0, X_2 = -1) = P(X_1 = 0 | C_1) \times P(X_2 = -1 | C_1) \times P(C_1) + P(X_1 = 0 | C_2) \times P(X_2 = -1 | C_2) \times P(C_2)$$

a priori - equally probable

Naive - Assuming that features are conditionally independent of each other

Insensitive to the number of training examples

Calculate which is more likely

$$1. P(Outcome|x) = P(X_1|Outcome) * P(X_2|Outcome) * P(Outcome)$$

$$2. P(Not Outcome|x) = P(X_1|Not Outcome) * P(X_2|Not Outcome) * P(Not Outcome)$$

Calculate probability outcome occurs

$$\#1 / (\#1 + \#2)$$

Perceptron classifier

$$x_0\theta_0 + x_1\theta_1 + x_2\theta_2$$

of the previous weight vector

first instance will compute based on initial vector

For -1, first point checks == -1, != 1 for the points after

Update weight = - (x × learning_rate) for the misclassified instance

For +1, first point checks != -1, == 1 for the points after

Update weight = + (x × learning_rate) for the misclassified instance

points are linearly separable if there are no longer any misclassified points

not misclassified: sign(x) function having -1 when x is ≤ 0

Supervised learning

Identifying objects

Output would be a continuous value

Given many emails, you want to **determine** if they are Spam or Non-Spam emails.

Given historical weather records, **predict** if tomorrow's weather will be sunny or rainy.

Unsupervised learning

No labels (no output) Kmeans

Input: Given a set of data points P , number of clusters k

1. Randomly pick k points from P as centers $c_j = 1..k$
2. Iterate (until max-iterations or c_j no longer changes up to some threshold)
 - Assign each point to nearest center: $y_i = \arg \min \| p_i - c_j \|^2$
 - c_j = Re-estimate each center as mean of points assigned to it

Given a set of news articles from many different news websites, find out what are the main topics covered.

From the user usage patterns on a website, figure out what different groups of users exist.

RL

Planning problems

choices that maximise the outcome

Find S

1. Initialise h to most specific h
2. If the positive instance differs, then replace with ?

Version space

List all hypothesis, remove any hypothesis that is inconsistent with any training example

Candidate Elimination

prefers positive examples over negatives as there can only be one maximally specific hypothesis in each iteration ? don't care

no value

Set $S_0 = \langle \emptyset \emptyset \emptyset \emptyset \emptyset \emptyset \rangle$

$G_0, G_1 = \langle \text{?????} \rangle$

For the negative example,

Create h from the specific hypothesis when it differs based on previous instance

Remove those that are inconsistent from the previous general hypotheses

Keep S to the previous instance S

For the positive example,

replace with ? when it differs

Cost functions

Logistic regression

Is binary classification

Need meta strategies for multi class and it chooses the model that has the highest confidence in predictions

Only gives values 0 and 1 so it is not a good result

$$J(\theta) = \frac{1}{m} \sum_{j=1}^m \ln(1 + \exp^{-y^j \theta^T x^j})$$

Linear regression

$$J(\theta) = \frac{1}{m} \sum_{j=1}^m (\theta^T x^j - y^j)^2$$

Linear algebra

$$\text{trace}(\text{AB}) = \text{trace}(\text{BA}) = \text{trace}(\text{I}_{n+1}) = n + 1$$

Idempotency

$$H^2 = H$$

$$(I - H)^2 = I - 2H + H = I - H$$

Pseudoinverse

$$(X^T X)^{-1} X^T$$

$$\theta = X^+ y$$

$$X = \begin{pmatrix} 1 & .. & .. \\ 1 & .. & .. \\ 1 & .. & .. \end{pmatrix}$$

Maximizing likelihood

$$\text{Minimize cross entropy} = - \frac{1}{m} \sum_{j=1}^m \ln g(y^j \theta^T x^j)$$

Ridge Regression

$$\frac{1}{m} \sum_{i=1}^m (y_i - w^T x_i)^2 + \lambda \|w\|^2$$

$$-2X^T(Y - X * \theta) + 2 * \alpha * \theta$$

Noise

Not deterministic only deterministic for $f(x) = 0$

$$y - f(x)$$

$$P(y - x)$$

$$\text{Noisy target} = E[y - x] + (y - f(x))$$

Stochastic noise vs Deterministic

To reduce stochastic noise, the only way is to re-measure y

To reduce deterministic noise, change H

Underfitting and Overfitting

Why do we regularize? We fit the data too much, we are fitting to the noise, balance fitting to the observed data. If we fit the data too well, the performance can be poor called overfitting. Overfitting can happen even if the data is not noisy.

Overfit - high level of noise, high complexity of $f(x)$

Underfit - high bias and low variance

Reduce bias by increasing complexity of hypotheses
simple representation of more complex reality

Non-linear transformation

Transform this non-linear equation to a linear one: $y = ax^b$?

Under the condition $y > 0$ and $x > 0$

$$\ln(y) = \ln(ax^b)$$

$$\ln(y) = \ln(a) + \ln(x^b)$$

$$\ln(y) = \ln(a) + b\ln(x)$$

let $y' = \ln(y)$, $x' = \ln(x)$, $a' = \ln(a)$, the equation can be written as
 $y' = a' + bx'$