

Some terms

Model-free: not learning the transition and reward, only need to get observations from the real world or simulator

Prediction to control - prediction into a control policy

Monte Carlo

$$\pi_{UCT}(n) = \operatorname{argmax}_a (Q(n, A) + c \sqrt{\frac{\log_2(N(n))}{N(n, a)}})$$

$$V(A) = \frac{\#A|A=1}{\#A + \#B}$$

Generalized policy iteration

Take argmax get the policy

Run the policy and estimate a new Q function

Don't have to evaluate the policy fully

Do one step and update one data (Q function)

Update greedily wrt the Q function

By being greedy (get a new policy) and improve it

Monte Carlo

Take many trajectories and average the trajectories

Temporal difference learning

Immediately update our prediction just after one step

But there is no return from the long trajectory

Current return of u

U is the estimate

Monte Carlo error: considers all the previous steps

vs.

TD error: how difference are my difference at this time step

n step TD

better tradeoff between bias and variance

When n is ∞ , it becomes Monte Carlo

Average can somewhat reduce variance (need less data to get good estimate,

learn faster as it converges faster)

TD(λ)

n step is weighted by $\frac{\lambda}{n-1}$

0 forget the future take 1 step

1 close to 1 important for a long time

We are averaging the estimates (more expensive)

ADP

Run v_i or p_i

not trying to estimate the model only the values

need to have simulator of the real world to go on with RL

Build a model using function approximation

What is the transition function to use?

What reward to use

Successful has been parameterised Q function, supervised learning etc.

Model-free are easier to work with

We usually mix MC with TD in practice, less estimate and not much variance

TD

Assumption is it is always markovian in real RL tasks in the environment

TD learning applies assumption to estimate

Initializer can be 0

A, 0, B, 0

Markovian assumption, if it sees A and B then it gives the reward of 0. Does not give the right values of the states because it makes an assumption that is not true of the environment

If it is true, then it will converge to correct value functions

Update equation

$$V^\pi(s) \leftarrow V^\pi(s) + \alpha(R(s) + \gamma V^\pi(s') - V^\pi(s))$$

$1 - \alpha$ can be obtained by extracting $-V^\pi(s)$ out and simplifying

SARSA

TD control

Uses GPI, on policy (improve the same policy)

target = $R(s) + \gamma Q(s', a')$

SARSA will converge to optimal policy on tabular cases

MC on the windy gridworld

Converges slower but eventually reaches optimal

Have to use old policy until end of the trajectory

SARSA changes policy at every time step, may wander around at the start, but will continue learning

ϵ greedy: Take the action suggested by the Q function or random exploration

Improving the current policy with ϵ greedy, when present does not converging to the optimal policy

trying to improve the policy

environment might change, world isn't quite a MDP, it can be robust, model may change over time

Might still fall off the cliff

Predictive problem learning (passive TD) has some restrictions

Update only at next state

Q-learning

Only have the q function (acts as policy on its own)

Not in SARSA: $\gamma \max_a$

best action at the next action of Q function

E.g currently have a deployed policy, at wild change the policy with SARSA, train the target policy

Behavior policy can be some other policy that is running, generate data to train the policy

If stop learning then Q-learning will outperform Sarsa

Similar concept to ϵ greedy

Let's go to part of the space that is exploit less

Issues when state space is large

Update policy then take argmax - given the current policy from the Q function, estimate a new Q function and construct a new policy by taking argmax of Q function

Q learning can learn off policy (doesn't follow the policy given, uses different policy)

Take the max over the next state

Minimizing the difference between the target and the current estimate

If the count is small give bonus, if we try many times reduce bonus

Why don't u update all the Q at each state?

Update is only done for all the state-action

Not the final Q-table, it will interact with the environment and update the table

What if the transitions are non-deterministic?

probability to move to a state

Even it is probability, it will move to that state, no longer a probability so take Q value of that state, expected Q learning

Function approximation

Number of states are huge, huge state space

Approximate (estimate) Q value with function approximator

Figure out what are the important features and extract out to define each state E.g. for a chess, queen is worth 5 points, pawn worth, knight worth

RL will learn the parameter vector Θ

Allows generalisation observations so that it can work on entire state space

Small fraction of the state

Applications: Go, Deep neural network, OpenAI5

MCT Search on simulator, for online search value function and combine with MCTS

Do learning with function approximation

Runs in the form of ridge regression

Online learning

One example, learn an example and update the weight vector

Update theta to improve estimate, step in the direction of the negative gradient

large buffer D, old sample does not mix with the recent one

$(s_t, a_t, r_{t+1}, s_{t+1})$ max over a without the k
learn and freeze it and use as target θ^-
 θ^- Q function from current - c steps

High profile achievement as one architecture can work across 50 atari games
Set a new target and change every c steps

Extract with a linear function of features
E.g. position of the ball
Need two frames to determine the direction we are going
provably good - how long would it take to go to epsilon